

MASTER THESIS

Thesis submitted in fulfillment of the requirements for the degree of Master of Science in Engineering at the University of Applied Sciences Technikum Wien - Degree Program Mechatronics/Robotics

SAGE: Multi object semantic aware guided exploration with persistent memory

By: Kevin Eppacher, BSc

Student Number: 2310331013

Supervisor: Simon Schwaiger, MSc

Vienna, December 11, 2025

This work was conducted in the context of the project “Stadt Wien Kompetenzteam für Drohnentechnik in der Fachhochschulausbildung” (project number MA23 35-02, financed by the Department MA23 for Economic Affairs, Labour and Statistics of the City of Vienna).

Funded by



Economic Affairs,
Labour and Statistics



Declaration

“As author and creator of this work to hand, I confirm with my signature knowledge of the relevant copyright regulations governed by higher education acts (see Urheberrechtsgesetz /Austrian copyright law as amended as well as the Statute on Studies Act Provisions / Examination Regulations of the UAS Technikum Wien as amended).

I hereby declare that I completed the present work independently and that any ideas, whether written by others or by myself, have been fully sourced and referenced. I am aware of any consequences I may face on the part of the degree program director if there should be evidence of missing autonomy and independence or evidence of any intent to fraudulently achieve a pass mark for this work (see Statute on Studies Act Provisions / Examination Regulations of the UAS Technikum Wien as amended).

I further declare that up to this date I have not published the work to hand nor have I presented it to another examination board in the same or similar form. I affirm that the version submitted matches the version in the upload tool.“

Vienna, December 11, 2025

Signature

Kurzfassung

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like "Huardest gefburn"? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Schlagworte: Keyword1, Keyword2, Keyword3, Keyword4

Abstract

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Keywords: Keyword1, Keyword2, Keyword3, Keyword4

Acknowledgements

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like "Huardest gefburn"? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Contents

1	Introduction	1
1.1	Problem Statement	1
1.2	Scientific Contribution	4
1.3	Thesis Structure	5
2	State of the Art	5
2.1	Geometric Exploration	6
2.2	Vision-Language-Guided Exploration	6
2.3	Map Reconstruction and Persistent Semantic Mapping	7
2.4	Object Detection and Promptable Models	8
3	Methods	8
3.1	System Overview	9
3.2	Semantic Frontier Exploration	9
3.3	Persistent Semantic 3D Mapping	11
3.4	Promptable Zero-Shot Detection	12
3.5	Fusion Strategy	13
3.6	Behavior Tree for Semantic-Guided Exploration	15
4	Implementation	15
4.1	Simulation Environment	16
4.2	Dataset	17
4.3	Used Software	17
4.4	Used Hardware	17
4.5	Evaluation Metrics	17
5	Discussion and Results	24
5.1	Experiment 1: Benchmarking on Matterport Scenes	24
5.2	Experiment 2: Impact of Exploration–Memory Weighting	24
5.3	Experiment 3: Sensitivity to Semantic Map Granularity	25
5.4	Experiment 4: Effect of Multi-Source Semantic Fusion	25
5.5	Experiment 5: System Efficiency and Real-World Validation	25
6	Summary and Outlook	25
Bibliography		26
List of Figures		28

List of Tables	29
List of source codes	30
A Appendix A	31
B Appendix B	32

1 Introduction

The introduction of Transformer-based architectures [1] has opened new opportunities for integrating high-level semantic reasoning with low-level geometric navigation in robotics. Traditional robotic exploration methods have primarily focused on mapping unknown environments using geometric cues, often neglecting the rich semantic information available in visual and linguistic modalities [2]. However, recent advances in large language models (LLMs) and vision-language models (VLMs) have enabled robots to interpret and act upon complex, open-ended instructions provided in natural language. These developments have catalyzed a shift toward semantic reasoning and zero-shot generalization, allowing robots to understand unseen concepts beyond fixed, task-specific datasets.

Consequently, new applications have emerged that require robots not only to explore and map their surroundings but also to reason about the semantic structure and relationships within them. In service robotics, for instance, a mobile agent may be instructed to locate a specific object based on high-level descriptions such as “*find the red chair in the living room*”, rather than relying on a limited set of predefined categories such as those from COCO [3]. Similarly, in search and rescue operations, robots may be tasked with locating missing persons based on vague or incomplete contextual information, such as the assumption that an individual might be found “*in the bathroom*”. In industrial inspection, autonomous agents must identify structural anomalies or specific components within unstructured and partially observable environments, while in warehouse automation, robots must locate items or storage units that may not be consistently labeled or fully visible.

Across these domains, the integration of semantic understanding with autonomous navigation is becoming increasingly essential. Robots must be capable of interpreting abstract human instructions, reasoning about both spatial and semantic context, and conducting efficient searches in dynamic, partially known environments. These challenges motivate the development of a unified framework that bridges geometric exploration with semantic scene understanding, enabling autonomous agents to perform open-vocabulary, goal-directed exploration guided by high-level semantic input.

1.1 Problem Statement

Geometric exploration techniques are widely used for mapping unknown environments, by identifying frontiers (2D or 3D for Drones), and navigating towards the frontier with the highest information gain, which can be determined with several algorithms. Especially useful are swarms of robots exploring large areas collaboratively. However, these methods are primarily designed for coverage and mapping tasks, rather than for goal-directed exploration based on semantic understanding. As a result, they often lead to inefficient search behavior when

the objective is to locate specific objects or regions of interest defined by high-level semantic criteria.

Therefore, instead of using only LiDAR or depth sensors for geometric mapping, incorporating RGB camera input allows for semantic reasoning about the environment. In table1 several recent works are summarized that leverage the use of either pretrained vision-language models or reinforcement learning based policies, guiding the robot towards areas that are more likely to contain the target object or region based on semantic cues.

Approach	Training Required	Real-Time	Semantic Reasoning Model
VLFM [4]	✗ (zero-shot)	✓	BLIP-2 + GroundingDINO + SAM
SemUtil [5]	✗ (training-free)	✓	Mask R-CNN + CLIP + BERT
ESC [6]	✗ (zero-shot)	✓	GLIP + DeBERTa / ChatGPT reasoning
LGX [7]	✗ (zero-shot)	✓	GPT-3 + GLIP + BLIP
CoW [8]	✗ (zero-shot)	✓	CLIP similarity scoring
ZSON [9]	✓ (RL pretraining)	✗	CLIP-based RL policy
PONI [10]	✓ (supervised)	✗	Learned potential-field network
PIRLNav [11]	✓ (BC + RL)	✗	DINO-based CNN-RNN policy

Table 1: Overview of Semantic Zero-Shot and Trained Exploration Approaches

ZSON [9], PONI [10], and PIRLNav [11] utilize deep reinforcement learning (DRL) or supervised learning to train exploration policies that can generalize to unseen objects. While these methods demonstrate promising results, they often require extensive offline training and may struggle to adapt to new environments or object categories not encountered during training. In contrast, VLFM [4], SemUtil [5], ESC [6], LGX [7], and CoW [8] leverage pretrained vision-language models to perform zero-shot semantic exploration without additional training. These approaches enable real-time decision-making based on semantic cues extracted from RGB images, allowing robots to navigate towards areas likely to contain the target object or region. ESC [6] and LGX [7] further incorporate large language models (LLMs) to enhance reasoning capabilities, enabling more sophisticated interpretation of high-level instructions, but at the cost of increased computational requirements and potential latency issues. Even if VLFM [4] real-time capable, it uses multiple heavy models (GroundingDINO, SAM, BLIP-2) that are difficult to run on mobile robots with limited compute resources, especially it runs on a 16 GB VRAM GPU (if all models are used at once). These approaches primarily focus on short-term exploration strategies, often neglecting the importance of persistent memory for long-term efficiency in multi-object search tasks and storing previously acquired semantic information within a unified world model. A global memory would allow the robot to recall previously explored areas and detected objects, reducing redundant exploration and improving overall search efficiency.

Current solutions:

Approach	Training Required	Real-Time	Memory Representation	Exploration Integration
OneMap	\times (zero-shot)	✓	2D probabilistic feature field	✓ (frontier-based)
ConceptGraphs	\times (pretrained models)	✗	3D scene graph	✓ (LLM-planner)
SemExp	✓ (RL + supervised)	✗	2D semantic occupancy map	✓ (learned policy)
GeFF	✓ (ScanNet pretrain)	✓	Implicit 3D feature field	✗ (passive)
RayFronts	\times (foundation model)	✓	Hybrid voxel + ray field	✓ (semantic frontier)
VLMaps	\times (pretrained LSeg/CLIP)	✗	2.5D open-vocab grid	✓ (frontier-compatible)
Pigeon	✓ (RLVR fine-tune)	✓	Point-of-Interest snapshot memory	✓ (reasoning-aware)

Table 2: Overview of Persistent or Memory-Based Semantic Mapping Approaches.

Despite notable progress, these methods often suffer from heavy GPU requirements, offline training procedures, or noisy semantic maps that reduce navigation reliability. Some further depend heavily on object detectors, which can introduce false positives in open-vocabulary settings.

Core Gaps in Existing Work:

Limitation	Example Works	Implication
No persistent memory	VLFM, CoW, LGX, ESC	No long-term fusion or recall; repeated exploration of known areas.
Offline training required	ZSON, PONI, PIRLNav, SemExp	Heavy RL/supervised training; poor adaptability to new scenes.
No balance between exploration and memory	OneMap, RayFronts, VLMaps	Either passive mapping or short-term exploration; inefficient search.
No zero-shot exploration	VLFM, CoW, LGX	Detect novel objects but fail to explore unseen regions strategically.
Premapping needed	ConceptGraphs, VLMaps, GeFF	Depend on pre-recorded data; not suited for online autonomy.
Limited robustness	PONI, SemExp, PIRLNav	Closed-set categories; fragile under real-world variation.
Low real-world applicability	ConceptGraphs, VLMaps	High GPU cost or simulation-only; not deployable on mobile robots.

Table 3: Identified Gaps in Existing Semantic Exploration Frameworks.

Derived Requirements from the State of the Art:

- Zero-shot frontier exploration framework.
- Persistent 3D semantic mapping with confidence-based fusion.
- Use of pretrained models only (no additional training required).
- Hybrid fusion strategy to balance exploration and memory.
- Multi-source detection fusion for robust object identification.
- GPU-efficient design for real-time deployment on mobile robots.
- Modular architecture for easy adaptation and extension.
- Independence from fixed, pre-trained object category sets.

1.2 Scientific Contribution

This work contributes to the state of the art by introducing a hybrid semantic exploration framework that integrates zero-shot semantic frontier scoring with persistent 3D scene representation, enabling autonomous search guided by open-vocabulary text queries. The system combines real-time semantic reasoning during exploration with a long-term spatial memory, allowing the robot to dynamically balance between discovering new information and exploiting previously acquired knowledge.

Unlike previous approaches that focus exclusively on either geometric frontiers or static semantic maps, the proposed framework continuously fuses information from multiple semantic sources to maintain a unified, confidence-based world model. Adaptive weighting enables the robot to adjust its behavior between exploration and exploitation according to the reliability of recent observations and the stability of stored semantic memory.

The framework further investigates how the quality and granularity of the underlying semantic information influence task success, navigation efficiency, and robustness. By systematically varying the trust between exploration and memory components, this work provides new insights into how semantic reasoning and persistent mapping can be effectively combined for open-vocabulary, multi-object search in dynamic environments.

To evaluate the contribution of the proposed system, the following research questions are formulated:

- 1. How does integrating zero-shot semantic exploration and persistent 3D semantic mapping affect multi-object search performance and navigation efficiency compared to existing methods?**

Metrics: Performance is quantified in terms of task success and path efficiency, measured through Success Rate (SR), Success per Path Length (SPL), and Multi-Object Success Rate (MSR) relative to representative state-of-the-art systems such as OneMap, VLFM, and Pigeon.

2. How does the interaction between live exploration and accumulated semantic memory influence overall system performance?

Metrics: The weighting factor between exploration and memory is varied during graph node fusion to assess impacts on SR and SPL, identifying optimal trade-offs between reactivity and exploitation.

3. How does the granularity of semantic map retrieval affect map quality, and can dynamic weighting between exploration and memory compensate for potential noise?

Metrics: The semantic granularity in the 3D semantic mapper is varied while adjusting exploration weight to evaluate effects on SR and SPL.

4. How does multi-source fusion of detection confidence, semantic similarity, and memory confidence impact detection robustness and false-positive suppression during exploration?

Metrics: Precision, Recall, F1-Score, Confusion Matrix, and SR under different fusion weight configurations across COCO, open-vocabulary, and zero-shot classes.

5. What is the computational footprint and real-world robustness of the hybrid framework?

Metrics: Frames per second (FPS), GPU/CPU usage, inference latency, and detection stability under sensor noise during physical deployment on a mobile robot.

These research questions guide the design of the experimental evaluation, where each question is systematically addressed through targeted ablation studies, comparative benchmarks, and real-world validation presented in Chapter ??.

1.3 Thesis Structure

2 State of the Art

In this chapter, the current state of research in semantic multi-object search, map reconstruction, and object detection is reviewed. The goal is to identify strengths and limitations of existing methods and establish the technological context for the proposed hybrid approach. The chapter is divided into three key areas: approaches for searching multiple objects semantically, techniques for building and maintaining persistent semantic maps, and recent advances in object detection and promptable models for open-vocabulary tasks.

2.1 Geometric Exploration

- Geometric exploration aims to explore unknown environments by navigating toward frontiers — the boundaries between known and unknown space — typically within a SLAM-based mapping framework.
- Selecting the frontier closest to the robot minimizes path cost (*greedy strategy*) [**topiwala2018frontierbasedexplorationautonomous**]. *Advantage:* simple and efficient. *Limitation:* ignores potential information gain.
- Prioritizing frontiers based on maximum expected information gain (e.g., entropy, mutual information) enables more informative exploration [**bourgault2002informationbasedadaptiveexploration**, **Suresh_2024**].
- Selecting the largest frontier clusters improves stability and continuity in exploration [**stachniss2009efficientexploration**].
- These approaches are purely geometric and lack semantic awareness. They are not suitable for open-vocabulary object queries or multi-object search tasks requiring semantic reasoning.

2.2 Vision-Language-Guided Exploration

- Review of methods targeting simultaneous or sequential search for multiple objects in unknown environments.
- Analysis of ...
 - VLFM [**yokoyama2024vlfm**]
 - SemUtil [**chen2023semutil**]
 - ESC [**liu2023esc**]
 - LGX [**shah2024lgx**]
 - CoW [**qi2023cow**]
 - ZSON [**shah2023zson**]
 - PONI [**qi2024poni**]
 - PIRLNav [**ma2024pirlnav**]

regarding aspects such as:

- Training Required (Pretrained vs. Fine-tuned)
- Real-Time Capability
- VRAM Requirements
- Sensor Modalities (RGB, RGB-D, LiDAR)
- Semantic reasoning (BLIP2, CLIP, GPT-4V, etc.)

- Persistent Memory (None, 2D Map, 3D Map, Scene Graph)
- Evaluation of performance metrics used in multi-object search:
 - Success Rate (SR)
 - Success weighted by Path Length (SPL)
 - MSR (Multi-object Success Rate)
- Discussion of semantic exploration frameworks combining language models with spatial reasoning.
- Challenges of maintaining semantic context across multiple targets.

2.3 Map Reconstruction and Persistent Semantic Mapping

- Overview of approaches to build and update semantic maps during exploration:
 - 2D grid maps
 - Pointclouds
 - Voxel grids
 - Octomaps
 - Scene graphs
 - Neural Radiance Fields (NeRFs)
 - Feature Fields
- Techniques for fusing sensor data into persistent 2D/3D representations:
 - Storing Visual Embeddings (e.g., CLIP features) in 3D maps for semantic querying.
 - Incremental updating of semantic labels based on new observations.
 - Handling uncertainty and conflicting detections over time.
- Comparison of representations (Octomaps, point clouds, voxel grids) in terms of:
 - Memory efficiency.
 - Ability to store semantic labels persistently.
- Discussion of ...
 - ConceptGraphs
 - ConceptGraph-Online
 - OpenFusion
 - Clio
 - OpenScene
 - GeFF

- CLIP-Fields
- ConceptFusion
- VLMaps
- LERF

as examples of global 3D semantic maps.

- Limitations in updating or correcting the map after wrong detections.

2.4 Object Detection and Promptable Models

- Review of traditional and open-vocabulary object detection methods.
- Analysis of grounding-capable detectors and segmentation models for zero-shot tasks.
- Specific evaluation of the following models for their suitability in semantic multi-object search:
 - YOLOv7
 - GroundingDINO
 - MobileSAM
 - GroundedSAM
 - SEEM
 - OWL-ViT
 - MaskDINO
- Discussion of promptable vision-language models supporting multi-modal queries (text, image, audio).
- Challenges with false positives in zero-shot settings and their implications for reliable multi-object detection.

3 Methods

This chapter details the methods developed for semantic exploration, persistent 3D mapping, promptable object detection, and robust fusion strategies for multi-object search.

3.1 System Overview

- Presentation of the overall architecture of the exploration (3.2), detection (3.4), mapping (3.3), and fusion (3.5) pipeline.
- Description of data flow between exploration (frontier evaluation), detection (promptable models), and memory (persistent semantic mapping), as shown in Figure 8.
- Explanation of how exploration and mapping components interact to progressively build a semantic understanding of the environment.




Figure 1: System architecture

3.2 Semantic Frontier Exploration

Exploration 2D Occupancy Map

- The SLAM map is used for navigation.
- Generating frontiers for each prompt would require deleting and rebuilding the SLAM map.
- This approach is inefficient and impractical for navigation.
- Therefore, a separate 2D occupancy grid is created exclusively for exploration.
- The exploration map is constructed by:
 - collecting robot poses, and
 - performing raytracing against the parent SLAM map.
- For each new semantic prompt:

- all stored poses are cleared, and
- the exploration occupancy grid is rebuilt from scratch.

Frontier Detection and Calculation

- Detection of frontiers on a 2D occupancy grid to identify candidate regions for exploration.
- Application of classical frontier-based exploration algorithms extended with semantic information.




Figure 2: Frontier detection on occupancy grid

Value Map Generation using Vision-Language Models

- Computation of value maps by evaluating cosine similarity between text queries and scene observations.




Figure 3: System architecture




Figure 4: Value map example

- Dynamic update of value maps as new observations are integrated.

3.3 Persistent Semantic 3D Mapping

Global Map Construction with Open-Fusion

- Incremental creation of a global semantic point cloud map integrating RGB-D observations over time.
- Registration of observations using robot poses to maintain a consistent world representation.
- Association of semantic labels with 3D points based on query relevance scores.

Semantic Clustering and Graph Node Generation

- Clustering of points with similar semantic labels to form object-level hypotheses.
- Construction of semantic graph nodes representing detected object instances with aggregated confidence scores.
- Maintenance of the semantic graph as a persistent memory for multi-object search tasks.

3.4 Promptable Zero-Shot Detection

- In this work YOLO-E [**yolo_e**] is used as the promptable zero-shot detection model.
- YOLO-E has the following advantages:
 - High inference speed suitable for real-time applications.
 - Ability to handle open-vocabulary object detection based on text prompts.
 - Integration of both visual and textual information for robust detection.
 - Pre-trained on large-scale datasets, enabling zero-shot generalization to unseen object categories.

Open-Vocabulary Object Detection with YOLO-E

- Utilization of the YOLO-E model for open-vocabulary object detection based on text prompts.
- Extraction of 2D bounding boxes and associated confidence scores for detected objects.
- Segmentation of detected objects to isolate relevant pixels for 3D localization.

Depth-Based 3D Localization

- With camera intrinsics and depth information, the 2D bounding boxes and segmentation masks are projected into 3D space.
- Calculation of 3D coordinates for each detected object using depth values within the bounding box.
- Semantic detection pointclouds are passed and then clustered and the centroid of each cluster is computed to obtain robust 3D object locations.
- For each cluster, the mean of the confidence scores of the associated 2D detections is calculated to assign a confidence score to the 3D localization.




Figure 5: YOLO-E detection to graph node 3D localization

3.5 Fusion Strategy

Exploration–Memory Weighting

- Exploration and memory graph nodes are fused and weighted as follows:
 - Proximity weighting: Nodes closer to the robot's current position are given higher weights.
 - Exploration vs Memory: Nodes from the exploration source are prioritized over memory nodes to encourage discovery of new information.
 - Costmap weighting: Nodes located in areas with lower navigation costs are favored to optimize path planning and navigation efficiency.

Multi-Source Detection Fusion

- Detection graph nodes are weighted based on:
 - YOLO-E confidence scores: Higher confidence detections are given more weight.
 - BLIP-2 value map: Detections with higher semantic relevance to the text prompt are prioritized.

3D Relevance Map – Combined Radial & Angular Gaussian




Figure 6: Fusion strategy for exploration, detection, and memory graph nodes

- The nearer detection graph nodes are to memory graph nodes, the higher their weight.

Relevance Filtering and Node Suppression

- Each source's graph nodes are filtered based on a relevance threshold to eliminate graph nodes within the fov map.
- Relevance map is build over time
- If a graph node is located in an area that has already been explored and found to be irrelevant to the prompt, it is suppressed.




Figure 7: Fusion strategy for exploration, detection, and memory graph nodes

3.6 Behavior Tree for Semantic-Guided Exploration

High-Level Task Structure

- The behavior tree (BT) is designed to manage the high-level task structure for semantic-guided exploration.
- The BT consists of the following main components:
 - Initialization: Clearing Maps, Publishing Prompts
 - Detection Branch: If object is detected over a threshold, navigate to it, realign to object take picture
 - Exploration Branch: While object not detected, perform semantic frontier exploration navigating to highest valued frontiers or memory nodes
 - Termination: If object found, end mission; If time limit reached, end mission
 - Behavior tree is called with a ros2 action server, which returns on termination success or failure, and actual path taken

Integration with Navigation Stack

- Navigation stack used for low-level path planning and obstacle avoidance.
- Action used: `navigate_to_pose`, `Spin`




Figure 8: System architecture

4 Implementation

This section details the practical implementation of the proposed approach, covering the simulation and real-world setup, datasets, software stack, and hardware configuration.

4.1 Simulation Environment

- Evaluation of simulation frameworks for indoor semantic navigation:
 - HabitatSim: Realistic Matterport3D-based environments with semantic annotations.
 - Isaac Sim / Isaac Lab: GPU-accelerated simulation, advanced physics, support for RTX ray tracing.
 - MuJoCo: High-speed physics engine, limited support for complex indoor scenes.
 - Ignition Gazebo: Modular simulator, ROS2 integration, good for real-robot transfer.

- ...

4.2 Dataset

- Use of **Matterport3D** scenes for realistic indoor environments with ground truth 3D reconstruction and semantic annotations.
- While the Habitat Navigation Challenge 2023 defines Success Rate (SR) and Success weighted by Path Length (SPL) as standard evaluation metrics within the Habitat-Sim environment, this work extends their application to Isaac Sim. Using Isaac Sim allows for a more physically accurate and sensor-consistent setup, incorporating realistic depth noise, lighting variation, and robot dynamics. To ensure comparability, SR and SPL are calculated following the official Habitat definitions, maintaining consistency with prior benchmarks while improving the realism of scene interaction and perception.

4.3 Used Software

- ROS2-based implementation (Humble Hawksbill) as middleware.
- Navigation stack: Navigation2 (Nav2) for frontier-based exploration and path planning.
- DDS communication layer for distributed communication between detection, mapping, and control nodes.
- Integration of promptable models (OpenFusion, BLIP-2, YOLO-E) for real-time zero-shot detection during exploration and exploitation.

4.4 Used Hardware

- **PC:**
 - CPU: AMD Ryzen 9 5950X 16-Core Processor
 - Motherboard: B550 Gaming X V2
 - GPU: ASUS TUF Gaming RTX 4090 24GB OC Edition
 - RAM: 64GB Corsair Vengeance LPX DDR4
- **Real Robot:** Configuration and components to be determined (TurtleBot Waffle).

4.5 Evaluation Metrics

This section defines the evaluation metrics used throughout the experiments and assigns them to each corresponding experiment.

Evaluation Pipeline Overview

- **Semantic Map Generation:** OpenFusion performs semantic segmentation of RGB-D input using the Matterport3D class list. Each segment is assigned its most likely class label from the detection model.
- **Manual Correction:** Incorrectly labeled segments can be manually relabeled within a dedicated ROS 2 node for semantic correction.
- **Data Storage:** OpenFusion saves both the 3D semantic pointcloud and the corresponding 2D SLAM map for each episode. All experiment data follows the `sage_datasets/matterport_isaac` directory structure.
- **Evaluation Initialization:** During evaluation, the saved maps and class definitions are loaded together with a list of target objects (e.g., “bed”, “toilet”, “chair”).
- **Class Filtering and Centroid Extraction:** The evaluator node filters the semantic pointcloud according to the target classes and extracts the 3D centroids of matching clusters.
- **Path Planning:** The shortest-path planner computes the geodesic-optimal path from the robot’s current pose to the nearest centroid of the selected target class, with the Global Path Planner from ROS2 Navigation2.
- **Metric Computation:** The evaluator node compares the planned and executed trajectories to compute Success Rate (SR), Success weighted by Path Length (SPL), and Multi-Object Success Rate (MSR).
- **Result Storage:** Evaluation metrics, trajectories, and intermediate results are stored per episode for analysis and benchmarking.

Experiment 1: Overall Performance Benchmarking

- Compare SR, SPL and MSR across different baseline methods and the proposed hybrid approach.
- Baselines include:
 - Vision-Language Frontier Maps for Zero-Shot Semantic Navigation (VLFM) [[shah2023vlm](#)],
 - Visual Language Maps for Robot Navigation (VLMaps) [[liu2023vlmaps](#)],
 - One Map to Find Them All (OneMap) [[yang20233don](#)],
 - **Pigeon! (Pigeon!)** [[qi2023pigeon](#)].
- Every scene within the ObjectNav HM3D v2 validation split:
- Within this dataset, each scene contains a set of episodes with the starting pose and target object specified.




Figure 9: Evaluation pipeline for benchmarking SR, SPL, and MSR in the semantic exploration framework.

- Due to the custom nature of the Isaac Sim environment, all baselines are re-implemented to ensure fair comparison under similar conditions.
- Limitations with IsaacSims Environment setup:
 - Starting Pose Variability
 - Amount of episodes per scene (5 per scene per floor with each 5 sub-episodes for multi-object search)
 - Different requirements:
 - * IsaacSim: Photorealistic rendering, physics simulation, realistic sensor noise
 - * HabitatSim: Optimized for fast navigation and large-scale datasets
- **Success Rate (SR):** Measures the proportion of tasks in which the robot successfully reaches the queried single goal object. This metric reflects the system's ability to se-

mantically ground a user-specified object and to navigate toward it reliably. It serves as a fundamental indicator of task success and is essential for evaluating overall system effectiveness in basic search scenarios. *Evaluation against:* VLFM, VLMaps, OneMap, Learning Generalizable Feature Fields for Mobile Manipulation (GeFF)

$$\text{SR} = \frac{1}{N} \sum_{i=1}^N S_i$$

where $S_i = 1$ if the goal was reached in episode i , and 0 otherwise; N is the total number of episodes.

- **Path Efficiency (SPL):** SPL measures the efficiency of successful navigation by comparing the shortest possible path to the actual path taken. It is defined only for successful runs and penalizes overly long trajectories. In the context of semantic exploration, SPL provides insight into how effectively the system prioritizes relevant regions and minimizes detours when searching for target objects.

$$\text{SPL} = \frac{1}{N} \sum_{i=1}^N S_i \cdot \frac{l_i}{\max(p_i, l_i)}$$

where S_i is the success indicator for episode i , l_i is the shortest path length to the goal, p_i is the actual path length taken, and N is the total number of episodes.

- **Multi-Object Success Rate (MSR):** The average number of successfully found objects per episode (*Progress, PR*) captures partial success in multi-goal navigation. SPL is computed separately for each object in sequence, conditioned on the success of the previous one. This highlights the system's ability to reuse semantic map information and improve efficiency across successive targets.

$$\text{PR} = \frac{1}{N} \sum_{i=1}^N C_i$$

where C_i is the number of successfully found objects in episode i , and N is the total number of episodes.

Experiment 2: Exploration–Memory Fusion Weighting

- **Objective:** Evaluate how varying the weighting between live semantic exploration and persistent 3D semantic memory influences navigation performance, stability, and overall search efficiency in the hybrid framework.
- **Fusion Parameter:** The trade-off between exploration and memory is controlled through a scalar weighting parameter

$$\lambda_{\text{exploration}} \in [0, 1],$$

which determines the relative influence of frontier-based exploration versus memory-driven exploitation during graph node fusion.

- **Research Questions:**

- **RQ2a:** How do Success Rate (SR) and Success weighted by Path Length (SPL) vary as the weighting shifts from exploration ($\lambda \rightarrow 1$) toward memory-driven behavior ($\lambda \rightarrow 0$)?
- **RQ2b:** Which weighting configuration yields the best trade-off between reactivity (fast adaptation to newly observed information) and stability (robust semantic localization using persistent memory)?

- **Evaluation Procedure:**

- Multiple runs are conducted across a range of $\lambda_{\text{exploration}}$ values.
- Performance is measured using the metrics SR and SPL.

- **Expected Outcome:** This experiment highlights whether hybrid fusion provides measurable benefits over purely exploration-driven or purely memory-driven behavior, and identifies the optimal balance for multi-object search tasks.

Experiment 3: Sensitivity to Semantic Map Granularity

- **Objective:** Investigate how the granularity of semantic retrieval in the 3D semantic mapper affects global map quality, navigation performance, and robustness of the hybrid exploration system. Specifically, this experiment evaluates whether dynamic rebalancing between exploration and memory can compensate for increased semantic noise introduced by higher retrieval depths.
- **Semantic Granularity Parameter:** Semantic map quality is controlled through the retrieval depth top-k, which specifies how many semantic candidates (from the VLM embedding space) are fused into each voxel:
 - Low top-k: sharper but potentially incomplete semantics.
 - High top-k: denser semantics but increased noise and ambiguity.
- **Dynamic Fusion Weighting:** To counteract noise introduced by larger top-k values, the exploration weight

$$\lambda_{\text{exploration}}$$

is progressively increased, shifting trust toward frontier-driven exploration and away from noisy memory components.

- **Research Questions:**

- **RQ4a:** How do Success Rate (SR) and Success weighted by Path Length (SPL) degrade as top-k increases while relying primarily on memory?
- **RQ4b:** Can adaptive rebalancing toward exploration (i.e., increasing $\lambda_{\text{exploration}}$) restore stable performance at higher top-k values?

- **Evaluation Metrics:**

- SR: ability to consistently reach goal objects under different semantic retrieval granularities.
- SPL: navigation efficiency and the influence of semantic noise on path quality.
- Additional qualitative assessment of map sharpness, cluster correctness, and temporal stability of semantic memory.

- **Evaluation Procedure:**

- Generate semantic maps at multiple top-k levels (e.g., 1, 3, 5, 10).
- For each top-k:
 - * Evaluate SR and SPL under memory-dominant settings.
 - * Incrementally increase $\lambda_{\text{exploration}}$ and re-evaluate.
- Compare results to determine:
 - * tolerance of the system to semantic noise,
 - * optimal balance between exploration and memory at different granularity levels,
 - * interaction effects between map resolution and fusion stability.
- **Purpose:** This experiment analyzes the coupling between semantic map granularity and the stability of exploration–memory fusion. Results reveal how coarse or noisy semantic retrieval affects overall navigation and whether adaptive weighting can maintain robust performance in open-vocabulary mapping environments.

Experiment 4: Robustness to False Positives Through Multi-Source Detection Fusion

- **Objective:** Evaluate how combining multiple semantic evidence sources, instance detection (YOLO-E), semantic similarity scoring (BLIP-2), and memory confidence from the 3D semantic map, improves detection robustness and suppresses false positives during exploration.
- **Fusion Model (Weighted Noisy-OR):** The proposed fusion strategy follows a weighted Noisy-OR formulation, in which independent semantic evidence sources jointly increase the probability of a valid detection:

$$S_{\text{fusion}} = 1 - (1 - w_d S_{\text{det}})(1 - w_c S_{\text{map}})(1 - w_m S_{\text{mem}}).$$

Here,

- S_{det} : YOLO-E detector confidence,
- S_{map} : similarity score from the value map (BLIP-2),
- S_{mem} : confidence from persistent 3D semantic memory,

- w_d, w_c, w_m : weights defining the contribution of each source.

This formulation ensures that high confidence from any source can compensate for uncertainty in others while suppressing spurious detections that lack multi-source agreement.

- **Research Questions:**

- **RQ3a:** How does overall task performance (SR) change under different weight configurations (w_d, w_c, w_m)?
- **RQ3b:** How do precision, recall, F1-score, and false-positive rate vary across:
 - * COCO-style closed-set categories,
 - * open-vocabulary object classes,
 - * zero-shot categories not seen during detector training?
- **RQ3c:** What drawbacks arise when detection thresholds are increased or when a single evidence source is overemphasised (e.g., memory bias, detector hallucination, missed low-confidence but valid detections)?
- **Evaluation Metrics:** Robustness is quantified using classification metrics derived from the confusion matrix:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN},$$

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}},$$

$$\text{FPR} = \frac{FP}{FP + TN}.$$

Additionally, downstream Success Rate (SR) is recorded for each weight triplet (w_d, w_c, w_m) to evaluate the effect of false positives on the overall navigation pipeline.

- **Evaluation Procedure:**

- Evaluate a range of weight combinations (w_d, w_c, w_m) spanning detector-dominant, map-dominant, memory-dominant, and balanced regimes, whereas the detection source is mandatory (i.e., $w_d > 0$).
- Compare against single-source baselines:
 - * detector-only (YOLO-E),
 - * similarity-only (BLIP-2),
 - * memory-only retrieval,
 - * the full Noisy-OR fusion strategy.
- Analyse outcomes under:
 - * closed-set (COCO) categories,
 - * open-vocabulary targets,
 - * zero-shot targets.

- Quantify how false positives propagate into:
 - * erroneous graph node generation,
 - * unnecessary navigation actions,
 - * degraded SR and SPL.
- **Purpose:** This experiment evaluates whether multi-source, Noisy-OR-based semantic fusion provides a measurable improvement in detection robustness and false-positive suppression compared to single-source methods, thereby enabling more reliable semantic exploration in open-vocabulary indoor environments.

Experiment 5: Real-World System Performance:

- SR, MSR and SP for search performance under real-world conditions.
- System metrics: CPU/GPU usage, FPS, inference latency.

Objective: Assess robustness, efficiency, and deployability in physical environments.

5 Discussion and Results

This chapter presents the experimental evaluation of the proposed hybrid semantic exploration system. Each experiment targets a specific research question and is evaluated using quantitative performance metrics.

5.1 Experiment 1: Benchmarking on Matterport Scenes

Evaluates baseline performance in multi-object search compared to state-of-the-art frameworks (OneMap, VLFM, Pigeon) using SR, SPL, and MSR.

5.2 Experiment 2: Impact of Exploration–Memory Weighting

Analyzes how varying the balance between live exploration and persistent memory influences navigation efficiency and task success.

5.3 Experiment 3: Sensitivity to Semantic Map Granularity

Investigates how varying the semantic retrieval depth affects mapping robustness and overall navigation stability.

5.4 Experiment 4: Effect of Multi-Source Semantic Fusion

Examines how combining detection confidence, semantic similarity, and memory reliability improves detection robustness and reduces false positives.

5.5 Experiment 5: System Efficiency and Real-World Validation

Assesses runtime performance, resource utilization, and stability under real-world sensor noise during physical deployment.

6 Summary and Outlook

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Bibliography

- [1] A. Vaswani *et al.* "Attention Is All You Need." Comment: 15 pages, 5 figures. arXiv: [1706.03762 \[cs\]](https://arxiv.org/abs/1706.03762). (Aug. 2, 2023), [Online]. Available: <http://arxiv.org/abs/1706.03762> (visited on 12/11/2025), pre-published.
- [2] A. Topiwala, P. Inani, and A. Kathpal. "Frontier Based Exploration for Autonomous Robot." arXiv: [1806.03581 \[cs\]](https://arxiv.org/abs/1806.03581). (Jun. 10, 2018), [Online]. Available: <http://arxiv.org/abs/1806.03581> (visited on 12/11/2025), pre-published.
- [3] T.-Y. Lin *et al.* "Microsoft COCO: Common Objects in Context." Comment: 1) updated annotation pipeline description and figures; 2) added new section describing datasets splits; 3) updated author list. arXiv: [1405.0312 \[cs\]](https://arxiv.org/abs/1405.0312). (Feb. 21, 2015), [Online]. Available: <http://arxiv.org/abs/1405.0312> (visited on 12/11/2025), pre-published.
- [4] N. Yokoyama, S. Ha, D. Batra, J. Wang, and B. Bucher. "VLFM: Vision-Language Frontier Maps for Zero-Shot Semantic Navigation." version 1. (2023), [Online]. Available: <https://arxiv.org/abs/2312.03275> (visited on 12/11/2025), pre-published.
- [5] J. Chen, G. Li, S. Kumar, B. Ghanem, and F. Yu. "How To Not Train Your Dragon: Training-free Embodied Object Goal Navigation with Semantic Frontiers." Comment: Accepted by/To be published in Robotics: Science and Systems (RSS) 2023; 11 pages, 5 figures. arXiv: [2305.16925 \[cs\]](https://arxiv.org/abs/2305.16925). (May 26, 2023), [Online]. Available: <http://arxiv.org/abs/2305.16925> (visited on 12/11/2025), pre-published.
- [6] K. Zhou *et al.* "ESC: Exploration with Soft Commonsense Constraints for Zero-shot Object Navigation." arXiv: [2301.13166 \[cs\]](https://arxiv.org/abs/2301.13166). (Jul. 6, 2023), [Online]. Available: <http://arxiv.org/abs/2301.13166> (visited on 12/11/2025), pre-published.
- [7] V. S. Dorbala, J. F. Mullen, and D. Manocha, "Can an Embodied Agent Find Your "Cat-shaped Mug"? LLM-Guided Exploration for Zero-Shot Object Navigation," *IEEE Robotics and Automation Letters*, vol. 9, no. 5, pp. 4083–4090, May 2024, Comment: 10 pages, ISSN: 2377-3766, 2377-3774. DOI: [10.1109/LRA.2023.3346800](https://doi.org/10.1109/LRA.2023.3346800). arXiv: [2303.03480 \[cs\]](https://arxiv.org/abs/2303.03480). [Online]. Available: <http://arxiv.org/abs/2303.03480> (visited on 12/11/2025).
- [8] S. Y. Gadre, M. Wortsman, G. Ilharco, L. Schmidt, and S. Song. "CoWs on Pasture: Baselines and Benchmarks for Language-Driven Zero-Shot Object Navigation." arXiv: [2203.10421 \[cs\]](https://arxiv.org/abs/2203.10421). (Dec. 14, 2022), [Online]. Available: <http://arxiv.org/abs/2203.10421> (visited on 12/11/2025), pre-published.
- [9] A. Majumdar, G. Aggarwal, B. Devnani, J. Hoffman, and D. Batra. "ZSON: Zero-Shot Object-Goal Navigation using Multimodal Goal Embeddings." Comment: code: <https://github.com/gunagg/zson>. arXiv: [2206.12403 \[cs\]](https://arxiv.org/abs/2206.12403). (Oct. 13, 2023), [Online]. Available: <http://arxiv.org/abs/2206.12403> (visited on 12/11/2025), pre-published.

- [10] S. K. Ramakrishnan, D. S. Chaplot, Z. Al-Halah, J. Malik, and K. Grauman. “PONI: Potential Functions for ObjectGoal Navigation with Interaction-free Learning.” Comment: 8 pages + supplementary. Accepted in CVPR 2022. arXiv: [2201.10029 \[cs\]](https://arxiv.org/abs/2201.10029). (Jun. 17, 2022), [Online]. Available: <http://arxiv.org/abs/2201.10029> (visited on 12/11/2025), pre-published.
- [11] R. Ramrakhya, D. Batra, E. Wijmans, and A. Das. “PIRLNav: Pretraining with Imitation and RL Finetuning for ObjectNav.” Comment: 8 pages + supplement. arXiv: [2301.07302 \[cs\]](https://arxiv.org/abs/2301.07302). (Mar. 26, 2023), [Online]. Available: <http://arxiv.org/abs/2301.07302> (visited on 12/11/2025), pre-published.

List of Figures

Figure 1 System architecture	9
Figure 2 Frontier detection on occupancy grid	10
Figure 3 System architecture	11
Figure 4 Value map example	11
Figure 5 YOLO-E detection to graph node 3D localization	13
Figure 6 Fusion strategy for exploration, detection, and memory graph nodes	14
Figure 7 Fusion strategy for exploration, detection, and memory graph nodes	15
Figure 8 System architecture	16
Figure 9 Evaluation pipeline for benchmarking SR, SPL, and MSR in the semantic exploration framework.	19

List of Tables

Table 1 Overview of Semantic Zero-Shot and Trained Exploration Approaches	2
Table 2 Overview of Persistent or Memory-Based Semantic Mapping Approaches.	3
Table 3 Identified Gaps in Existing Semantic Exploration Frameworks.	3

List of source codes

A Appendix A

B Appendix B